* Completed the first major maintenance window 12-15 hour day
* Fixed network configuration consistency across an AI fleet.
* Standardized hostnames by GPU type and BMC/XCC access.
* Standardized DHCP/netplan behavior
* Isolated rack focused for maintenance from the power spaghetti
* Re-cabled and cleaned the entire rack.
* Improved airflow.
* Reduced operational risk.
* Verified inventory
* Created actual documentation evidence with before/after photos.
* Kicked off multi-terabyte migrations.
* Solved the BTRFS mounting problem.
* Evaluated workstation strategy.
* Evaluated storage architecture strategy.
* Got management visibility
* Got maintenance authority enforced

// fixing the inherited reality

* accumulated expediency
* everybody plugging things in
* no lifecycle management ownership
* airflow impedence as an afterthought

// Identified:

* rack too shallow
* PDU placement wrong
* power constraints
* mixed ownership
* missing blanks
* missing cable management accessories

// Scope of work

> Capacity planning.
> Resource allocating.
> Hardware lifecycle management.
> Migration strategy.
> Disaster recovery.

// Team communication post maintenance window

6/5 8:52 PM

Team,

This first maintenance is now complete, it's been a long day and extended configurations for everything.

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

Please update your records for the systems:

hostname changes to reflect gpu for faster identication
ip changes due to dhcp to specified ports

[redacted]

network configuration for os and for system bmc/xcc has been modified to be consistent across the fleet of systems

backend cabling of [redacted]-h200, [redacted]-h200, [redacted]-mi300x (all rack 1) have been recabled to remedy the thermal/airflow issue, in addition to detached from the others from an intertwined power stance

PLEASE DO NOT CONNECT TO:

[redacted]-h200 /[redacted]
[redacted]-h200 / [redacted]
as they are still doing data migration efforts

[redacted]-b200 / [redacted] has finished it's data migration

Once the data migration is finished/verified for :
[redacted]-h200 / [redacted]
[redacted]-h200 / [redacted]
[redacted]-b200 / [redacted]

I still need to reconfigure the systems to prepare for the requested drives removal.

The other systems should be available.
Please log into those systems to verify and let me know if you have concerns.

I will send a follow-up update to when
[redacted]-h200 / [redacted]
[redacted]-h200 / [redacted]
[redacted]-b200 / [redacted]
are available for use again

===

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

6/8 8:06 AM

Good Morning team,

As mentioned on Friday evening (US est),
PLEASE DO NOT CONNECT TO:

[redacted]-h200 / [redacted]
[redacted]-h200 / [redacted]
as they are still doing data migration efforts

[redacted]-b200 /[redacted] has finished it's data migration

===

Data migration efforts are in progress, if you connected, please stop the workload and log off, so the data migration efforts can continue, else whatever data migration efforts going on will be incomplete/corrupted.

It seems
[redacted]-h200 / [redacted]
[redacted]-h200 / [redacted]
have workloads active.

These migrations needs to be complete, so that the respective users can verify/validate their data before Wednesday 6/10 eod (US est) so that I can prepare the system reconfiguration to clean up the former drive shares to remove them and return them

If you are using
[redacted]-b200 / [redacted]
Please login to verify your data, in the new data-migration share and let me know if you have issues before Wednesday 6/10.

6/10 10:31 AM
Team,

The data-migration efforts should be done. Users of the affected systems ([redacted]) have been reached out to have them verify/validate the data to address any issues before this Thursday 6/11/26 @4pm (US est) as the 4x 30tb drives will be wiped, system reconfigured, and drives pulled to return to the equipment owner on Friday 6/12.

Reminder for users of [redacted]) and [redacted]
notification has been sent to complete any workloads by Friday 6/12/2026 @ 4pm (US est), the system will be powered down as the equipment owner has requested to have the pcie network adapters returned. Request has been deferred to early next week.

6/12 3:29 PM
Team,

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

Systems ([redacted]-h200, [redacted]-h200, [redacted]-b200) had the 4x 30tb drives wiped, system reconfigured, and drives pulled and returned to the equipment owner through proxy (handed off to [redacted])

[redacted]-h200 and [redacted]-h200 are back online and available for use.

I will be shutting down systems ([redacted]-b200 and [redacted]-mi300x) on Monday 6/15/2026 morning by 9am US est at latest, and let the equipment owner know they can come pull those pcie network adapters. I'll send out a follow-up when they are done and I've powered on the systems.

6/15 1:05 PM
Team,

Systems ([redacted]-b200 and [redacted]-mi300x) are back online and available. Equipment owner has reclaimed the pcie network adapters.

6/22 11:52 AM
Team,

I've been informed that you have a testing cycle coming up, and it's recommended to do a short notice maintenance, planned for this Friday 6/26/2026 @ 12 noon US est. The work will be to decouple the power connections from the 5 systems :

[redacted]-h200 / [redacted]
[redacted]-h100 / [redacted]
[redacted]-b200 / [redacted]
[redacted]-b200 / [redacted]
[redacted]-mi300x / [redacted]

This is to ensure that future maintenance on any one of those systems is just isolated to that specific system and not affecting the other intertwined systems.

Please voice any issues, concerns, questions, objections before end of day 6/24/2026 Wednesday @4pm US est. Else maintenance will proceed as scheduled, the systems will be offline until maintenance is complete and communicated as such.

Please ensure all workloads are completed and do not start any new workloads after 6/25/2026 Thursday @ 4pm.
Thanks

kt

6/22 4:00 PM

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

Team,

[redacted]-mi300x / [redacted] : storage reconfiguration
This system is getting its storage drives reconfigured
Users of this system have also been notified via email to

BACK UP AND VALIDATE YOUR DATA
Please do so by Friday 6/26/2026 by 4pm US est, after this the system storage capacity will be
reconfigured for data2, all data destroyed and drives pulled and repurposed in another system.

thanks,

kt

6/26 12:02 PM
Team,

Maintenance is starting to decouple the power connections from the 5 systems :

[redacted]-h200 / [redacted]
[redacted]-h100 / [redacted]
[redacted]-b200 / [redacted]
[redacted]-b200 / [redacted]
[redacted]-mi300x / [redacted]

This is to ensure that future maintenance on any one of those systems is just isolated to that specific system and not affecting the other intertwined systems.

[redacted]-mi300x / [redacted] : storage reconfiguration
This system is getting its storage drives reconfigured and excess drives repurposed to [redacted]-h100 / [redacted]

Affected systems will be down until maintenance completion is communicated again.

6/26 5:43 PM
Team,

The maintenance to decouple the power connections from the 5 systems :

[redacted]-h200 / [redacted]
[redacted]-h100 / [redacted]
[redacted]-b200 / [redacted]
[redacted]-b200 / [redacted]
[redacted]-mi300x / [redacted]

is complete.

//note the ip change for the following systems

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
[redacted]-h100 > [redacted]

Storage reconfiguration of [redacted]-mi300x / [redacted] is complete.
Storage drives reconfigured repurposed to [redacted]-h100 / [redacted]

Systems should be back up, please login to verify your data.
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
//maintenance started as scheduled with no objections raised during window of
opportunity
//maintenance executing as scheduled for 6/26/2026 at 12pm US est
//users logged in with active workloads


  IPv4 address for ens11f0np0: [redacted]
  IPv6 address for ens11f0np0: fd55:faaf:e1ab:2021:ba3f:d2ff:fe90:addc

  => / is using 91.1% of 195.80GB



[redacted]-h200:~$
 16:14:42 up 20 days, 19:27, 10 users,  load average: 20.68, 20.22, 21.63
USER     TTY      FROM             LOGIN@   IDLE   JCPU   PCPU WHAT
[redacted] pts/21   tmux(47088).%6   10Jun26  7:34m  0.08s  0.08s -bash
[redacted] pts/24   tmux(47088).%0   08Jun26 17days  1.42s  1.42s -bash
[redacted] pts/16   tmux(47088).%1   10Jun26  2:24m 58.28s 33.50s /data2/
[redacted]/miniconda3/envs/gemma4_qkl/bin/python3.11 /data2
[redacted] pts/17   tmux(47088).%2   10Jun26 32:14m  0.11s  0.11s -bash
[redacted] pts/18   tmux(47088).%3   10Jun26 16days  0.02s  0.02s -bash
[redacted] pts/19   tmux(47088).%4   10Jun26  4days  0.07s  0.07s -bash
[redacted] pts/20   tmux(47088).%5   10Jun26  7:34m  0.29s  0.29s -bash
[redacted] pts/6    tmux(47088).%7   Thu02   37:53m  0.04s  0.04s -bash
[redacted] pts/7    tmux(47088).%8   Thu02   37:52m  0.04s  0.04s -bash


====


  IPv4 address for ens11f0: [redacted]
  IPv6 address for ens11f0: fd55:faaf:e1ab:2021:8616:cff:feac:7462

  => / is using 92.2% of 294.23GB
  => /data is using 92.5% of 13.97TB

[redacted]-h100:~$



====


//maintenance started as scheduled with no objections raised during window of
opportunity
//maintenance executing as scheduled for 6/26/2026 at 12pm US est
//users logged in with active workloads


  IPv4 address for enp90s0f0np0: [redacted]
  IPv6 address for enp90s0f0np0: fd55:faaf:e1ab:2021:ba3f:d2ff:fe4c:dbc

  => /data2 is using 90.7% of 13.86TB

[redacted]-b200:~$
 16:16:54 up 9 days, 20:47,  3 users,  load average: 6.08, 8.22, 9.29
```

Personal Intellectual Property. Corporate Use Without Permission. ©Kim Tran

```
USER     TTY      FROM             LOGIN@   IDLE   JCPU   PCPU WHAT

[redacted] pts/4  tmux(1389067).%1 Thu07   18:52m  0.01s  0.01s -bash
         pts/6    tmux(1389067).%0 Wed02    2days  0.00s  0.00s -bash
@llmus03-b200:~$
```

====

//maintenance started as scheduled with no objections raised during window of opportunity
//maintenance executing as scheduled for 6/26/2026 at 12pm US est
//users logged in with active workloads

```
  IPv4 address for ens10f0: [redacted]
  IPv6 address for ens10f0: fd55:faaf:e1ab:2021:a236:9fff:feb2:d8dc

[redacted]-b200:~$ w
 16:18:07 up 2 days,  8:24,  9 users,  load average: 12.64, 10.65, 9.75
USER     TTY      FROM             LOGIN@   IDLE   JCPU   PCPU  WHAT

[redacted]        [redacted]       07:26    9:42m  0.00s  0.04s sshd: shuhq1 [priv]
                                   07:25    9:42m  0.00s  0.04s sshd: shuhq1 [priv]
                                   07:16    9:42m  0.00s  0.12s sshd: shuhq1 [priv]
                                   05:49    9.42m  0.00s  0.11s sshd: shuhq1 [priv]
-b200:~$
```

====

```
  IPv4 address for enp90s0f0np0: [redacted]
  IPv6 address for enp90s0f0np0: fd55:faaf:e1ab:2021:ac0:ebff:feff:379a

[redacted]-mi300x:~$
```

====

//system to be reconfigured for removal of /dev/md1 (data2)
//drives to be repurposed to system [redacted]-h100

```
  IPv4 address for ens11f0: [redacted]
  IPv6 address for ens11f0: fd55:faaf:e1ab:2021:6efe:54ff:fe3d:5138

  => / is using 85.2% of 435.96GB


[redacted]-mi300x:~$
Filesystem                          Size  Used Avail Use% Mounted on

/dev/md0                             11T  2.4T  8.2T  23% /data1
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
/dev/md1                       11T   75G   11T   1% /data2
[redacted]-mi300x:~$
```

```
[redacted]-mi300x:~$/data1
total 12K
4.0K drwxrwxr-x 11 [redacted] 4.0K Jun 12 02:02 [redacted]3
4.0K drwxrwxr-x 16 [redacted] 4.0K May 28 07:56 [redacted]g
   0 drwx--x--- 12 root        root        217 Jun 15 15:50 docker
   0 drwxrwxr-x  2 [redacted]  61 May 28 07:55 [redacted]
   0 -rwxrwxrwx  1 [redacted]   0 Apr 28 13:45 testfile
4.0K drwxrwxr-x 30 [redacted] 4.0K Jun 26 16:04 [redacted]
[redacted]-mi300x:~$
```

```
[redacted]-mi300x:~$/data2
total 0
[redacted]-mi300x:~$
```

```
[redacted]-mi300x:~$
/dev/md1:
           Version : 1.2
     Creation Time : Tue Apr 28 00:54:21 2026
        Raid Level : raid6
        Array Size : 11251445760 (10.48 TiB 11.52 TB)
     Used Dev Size : 1875240960 (1788.37 GiB 1920.25 GB)
      Raid Devices : 8
     Total Devices : 8
       Persistence : Superblock is persistent

     Intent Bitmap : Internal

       Update Time : Fri Jun 26 18:28:56 2026
             State : clean
    Active Devices : 8
   Working Devices : 8
    Failed Devices : 0
     Spare Devices : 0

            Layout : left-symmetric
        Chunk Size : 512K

Consistency Policy : bitmap

              Name : [redacted]:1
              UUID : bdf74c61:7540ae0a:7e803c78:fd0b58a0
            Events : 1793

    Number   Major   Minor   RaidDevice State
       0     259       20        0      active sync   /dev/nvme8n1p1
       1     259       30        1      active sync   /dev/nvme9n1p1
       2     259       26        2      active sync   /dev/nvme10n1p1
       3     259       34        3      active sync   /dev/nvme11n1p1
       4     259       43        4      active sync   /dev/nvme12n1p1
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
      5     259       42        5      active sync   /dev/nvme13n1p1
      6     259       47        6      active sync   /dev/nvme14n1p1
      7     259       44        7      active sync   /dev/nvme15n1p1
[redacted]mi300x:~$
```

```
[redacted]-mi300x:~$
/dev/md0:
           Version : 1.2
     Creation Time : Tue Apr 28 00:53:14 2026
        Raid Level : raid6
        Array Size : 11251445760 (10.48 TiB 11.52 TB)
     Used Dev Size : 1875240960 (1788.37 GiB 1920.25 GB)
      Raid Devices : 8
     Total Devices : 8
       Persistence : Superblock is persistent

     Intent Bitmap : Internal

       Update Time : Fri Jun 26 18:28:56 2026
             State : clean
    Active Devices : 8
   Working Devices : 8
    Failed Devices : 0
     Spare Devices : 0

            Layout : left-symmetric
        Chunk Size : 512K

Consistency Policy : bitmap

              Name : [redacted]:0
              UUID : dc6a9da3:aca8d91a:247568ed:71042b22
            Events : 1814

    Number   Major   Minor   RaidDevice State
       0     259        2        0      active sync   /dev/nvme0n1p1
       1     259        5        1      active sync   /dev/nvme1n1p1
       2     259       11        2      active sync   /dev/nvme2n1p1
       3     259       14        3      active sync   /dev/nvme3n1p1
       4     259       17        4      active sync   /dev/nvme4n1p1
       5     259       13        5      active sync   /dev/nvme5n1p1
       6     259       23        6      active sync   /dev/nvme6n1p1
       7     259       35        7      active sync   /dev/nvme7n1p1
[redacted]-mi300x:~$
```

```
[redacted]mi300x:~$
/dev/md0 on /data1 type xfs
(rw,relatime,attr2,inode64,logbufs=8,logbsize=32k,sunit=1024,swidth=6144,noquota)
/dev/md1 on /data2 type xfs
(rw,relatime,attr2,inode64,logbufs=8,logbsize=32k,sunit=1024,swidth=6144,noquota)
[redacted]mi300x:~$
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
[redacted]-mi300x:~$
# /etc/fstab: static file system information.
#
# Use 'blkid' to print the universally unique identifier for a
# device; this may be used with UUID= as a more robust way to name devices
# that works even if disks are added and removed. See fstab(5).
#
# <file system> <mount point>   <type>  <options>       <dump>  <pass>
# / was on /dev/ubuntu-vg/lv-0 during curtin installation
/dev/disk/by-id/dm-uuid-LVM-
UIH6kHQX1T38uhtmqrBi1Gpng3ycYCOr76mX0RItYKvSkcf5wuieJtygyoePpwq8 / ext4 defaults 0
1
# /boot was on /dev/sda2 during curtin installation
/dev/disk/by-uuid/fb6bca4d-0467-4ed0-95fa-6324c0aaa653 /boot ext4 defaults 0 1
# /boot/efi was on /dev/sda1 during curtin installation
/dev/disk/by-uuid/3353-AF76 /boot/efi vfat defaults 0 1
# manual entry for the xfs vols of md0
/dev/disk/by-uuid/aa02a646-11a2-4728-b13f-43daea9be1f3  /data1  xfs
defaults,nofail 0 2
# manual entry for the xfs vols of md1
/dev/disk/by-uuid/6f1ee2f2-d6cc-48d1-beb1-7859735f78a0  /data2  xfs
defaults,nofail 0 2
/swap.img       none    swap    sw      0       0
[redacted]-mi300x:~$
```

```
[redacted]-mi300x:~$
Personalities : [raid6] [raid5] [raid4] [linear] [multipath] [raid0] [raid1]
[raid10]
md1 : active raid6 nvme13n1p1[5] nvme12n1p1[4] nvme8n1p1[0] nvme14n1p1[6]
nvme15n1p1[7] nvme10n1p1[2] nvme11n1p1[3] nvme9n1p1[1]
      11251445760 blocks super 1.2 level 6, 512k chunk, algorithm 2 [8/8]
[UUUUUUUU]
      bitmap: 0/14 pages [0KB], 65536KB chunk

md0 : active raid6 nvme2n1p1[2] nvme5n1p1[5] nvme4n1p1[4] nvme3n1p1[3] nvme0n1p1[0]
nvme6n1p1[6] nvme1n1p1[1] nvme7n1p1[7]
      11251445760 blocks super 1.2 level 6, 512k chunk, algorithm 2 [8/8]
[UUUUUUUU]
      bitmap: 0/14 pages [0KB], 65536KB chunk

unused devices: <none>
[redacted]-mi300x:~$
```

```
[redacted]-mi300x:~$
# mdadm.conf
#
# !NB! Run update-initramfs -u after updating this file.
# !NB! This will ensure that initramfs has an uptodate copy.
#
# Please refer to mdadm.conf(5) for information about this file.
#

# by default (built-in), scan all partitions (/proc/partitions) and all
# containers for MD superblocks. alternatively, specify devices to scan, using
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
# wildcards if desired.
#DEVICE partitions containers

# automatically tag new arrays as belonging to the local system
HOMEHOST <system>

# instruct the monitoring daemon where to send mail alerts
MAILADDR root

# definitions of existing MD arrays

# This configuration was auto-generated on Wed, 11 Sep 2024 14:22:21 +0000 by
mkconf
ARRAY /dev/md0 metadata=1.2 name=llmus08:0 UUID=dc6a9da3:aca8d91a:247568ed:71042b22
ARRAY /dev/md1 metadata=1.2 name=llmus08:1 UUID=bdf74c61:7540ae0a:7e803c78:fd0b58a0
[redacted]-mi300x:~$



[redacted]-mi300x:~$
umount: /data2 unmounted
@llmus08-mi300x:~$
Filesystem                    Size  Used Avail Use% Mounted on

/dev/md0                       11T  2.4T  8.2T  23% /data1
[redacted]-mi300x:~$




[redacted]-mi300x:~$

/dev/md0 on /data1 type xfs
(rw,relatime,attr2,inode64,logbufs=8,logbsize=32k,sunit=1024,swidth=6144,noquota)
[redacted]-mi300x:~$



[redacted]-mi300x:~$
mdadm: stopped /dev/md0
[redacted]-mi300x:~$


[redacted]-mi300x:~$
Personalities : [raid6] [raid5] [raid4] [linear] [multipath] [raid0] [raid1]
[raid10]
md0 : active raid6 nvme2n1p1[2] nvme5n1p1[5] nvme4n1p1[4] nvme3n1p1[3] nvme0n1p1[0]
nvme6n1p1[6] nvme1n1p1[1] nvme7n1p1[7]
      11251445760 blocks super 1.2 level 6, 512k chunk, algorithm 2 [8/8]
[UUUUUUUU]
      bitmap: 0/14 pages [0KB], 65536KB chunk

unused devices: <none>
[redacted]-mi300x:~$



====
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
//upgraded raid pool to zfs raid6 for better data integrity and resiliency
[redacted]-h100:~$
data1 on /data1 type zfs (rw,xattr,noacl)
[redacted]-h100:~$


[redacted]-h100:~$ /data
total 8.0K
4.0K drwxrwxr-x 20 [redacted]     4.0K Apr 24 00:57 [redacted]
   0 drwxrwxr-x 14 [redacted]      308 Jun 25 20:37 [redacted]
4.0K drwxrwx--- 22 [redacted] 4.0K Jun 24 13:39 [redacted]
   0 drwxrwxr-x  5 [redacted]       69 Mar 12 16:17 [redacted]
[redacted]-h100:~$


[redacted]-h100:~$ /data1
total 512
512 -rw-rw-r-- 1  0 Jun 26 20:17 test-file
[redacted]-h100:~$


[redacted]-h100:~$
NAME     USED   AVAIL      REFER   MOUNTPOINT
data1    296K   10.3T      47.1K   /data1
[redacted]-h100:~$


[redacted]-h100:~$
  pool: data1
 state: ONLINE
config:

        NAME         STATE     READ  WRITE  CKSUM
        data1        ONLINE       0      0      0
          raidz2-0   ONLINE       0      0      0
            nvme2n1  ONLINE       0      0      0
            nvme3n1  ONLINE       0      0      0
            nvme0n1  ONLINE       0      0      0
            nvme1n1  ONLINE       0      0      0
            nvme6n1  ONLINE       0      0      0
            nvme7n1  ONLINE       0      0      0
            nvme4n1  ONLINE       0      0      0
            nvme5n1  ONLINE       0      0      0

errors: No known data errors
[redacted]-h100:~$


[redacted]-h100:~$
Filesystem                         Size   Used Avail Use% Mounted on

/dev/nvme8n1p1                      14T    13T  1.1T  93% /data
data1                               11T   128K   11T   1% /data1
[redacted]-h100:~$
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
====

//note the ip change for the following systems

[redacted]-h100 > [redacted]



====
```

systems verified back up:

```
$ ssh @[redacted]

  IPv4 address for ens11f0np0: [redacted]
  IPv6 address for ens11f0np0: fd55:faaf:e1ab:2021:ba3f:d2ff:fe90:addc

  => / is using 90.8% of 195.80GB

Last login: Fri Jun 26 20:52:56 2026
[redacted]-h200:~$



$ ssh [redacted]

  IPv4 address for ens11f0: [redacted]
  IPv6 address for ens11f0: fd55:faaf:e1ab:2021:8616:cff:feac:7462

  => / is using 92.0% of 294.23GB
  => /data is using 92.5% of 13.97TB

Last login: Fri Jun 26 20:55:20 2026
[redacted]-h100:~$



$ ssh @[redacted]

  IPv4 address for enp90s0f0np0: [redacted]
  IPv6 address for enp90s0f0np0: fd55:faaf:e1ab:2021:ba3f:d2ff:fe4c:dbc

  => /data2 is using 91.3% of 13.86TB

Last login: Fri Jun 26 20:58:36 2026
[redacted]-b200:~$




$ ssh @[redacted]

  IPv4 address for ens10f0: [redacted]
  IPv6 address for ens10f0: fd55:faaf:e1ab:2021:a236:9fff:feb2:d8dc

Last login: Fri Jun 26 19:50:30 2026 from [redacted]
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

```
[redacted]-b200:~$
```

```
$ ssh @[redacted]

  IPv4 address for enp90s0f0np0: [redacted]
  IPv6 address for enp90s0f0np0: fd55:faaf:e1ab:2021:ac0:ebff:feff:379a

Last login: Fri Jun 26 21:22:17 2026 from [redacted]
[redacted]-mi300x:~$
```

```
$ ssh @[redacted]

  IPv4 address for ens11f0: [redacted]
  IPv6 address for ens11f0: fd55:faaf:e1ab:2021:6efe:54ff:fe3d:5138

  => / is using 85.2% of 435.96GB

Last login: Fri Jun 26 18:27:59 2026 from [redacted]
[redacted]-mi300x:~$
```

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

My inheritance:

What I had to clean up and redeploy or remedy as best effort given the constraints and authorization

Initial fleet of 10, expanding into 13~16, each 8 gpu cluster system ~$1M cost/value

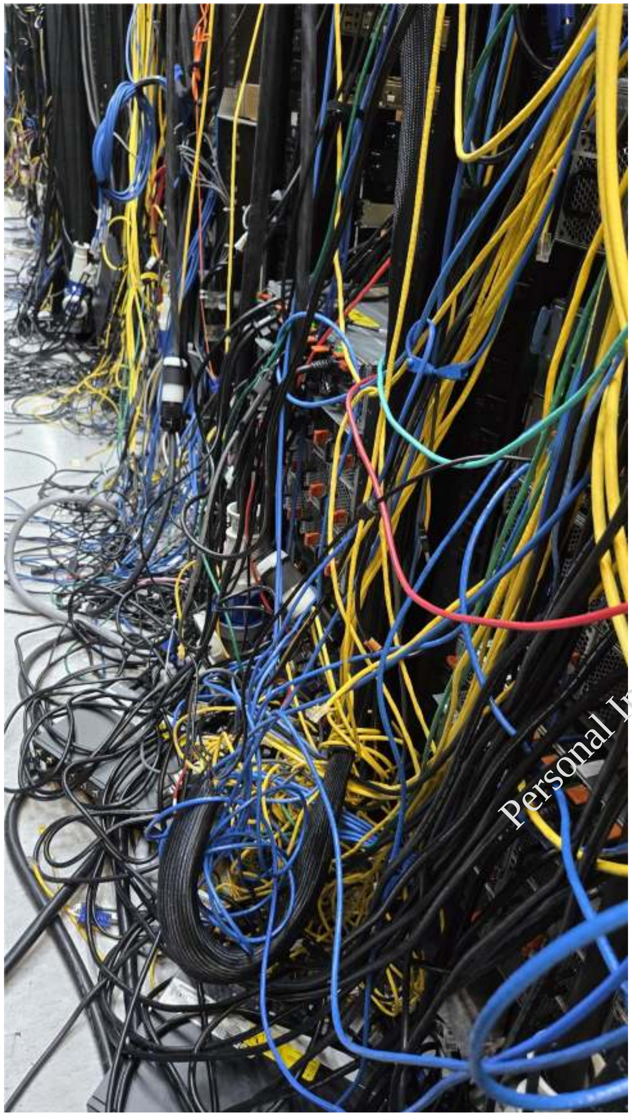

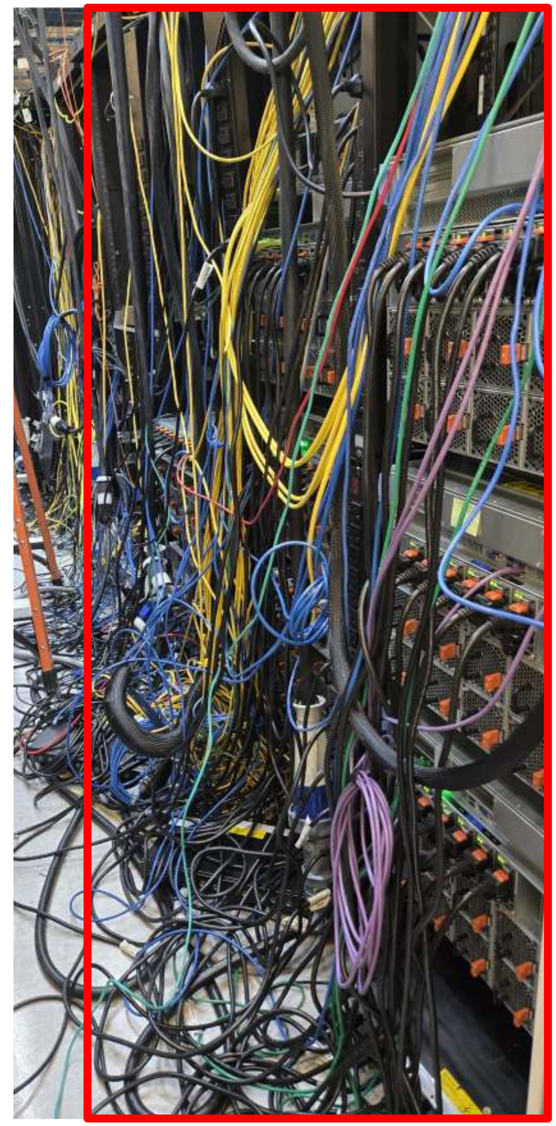

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

Thermal issues from the back and front, hot zones, no clearance for the engineered airflow

The spaghetti cabling, nothing traceable, no labels, no standard best practices of cabling lengths, added troubleshooting due to unable to rule out any known good gear (cabling)

The first 3 racks before any clean up attempts...

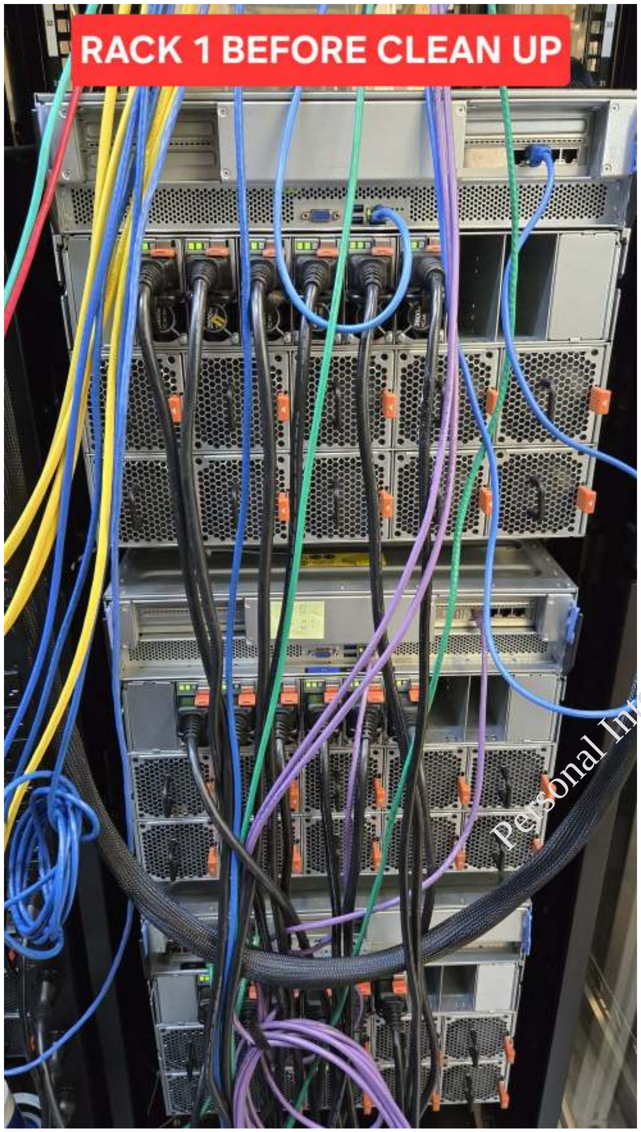

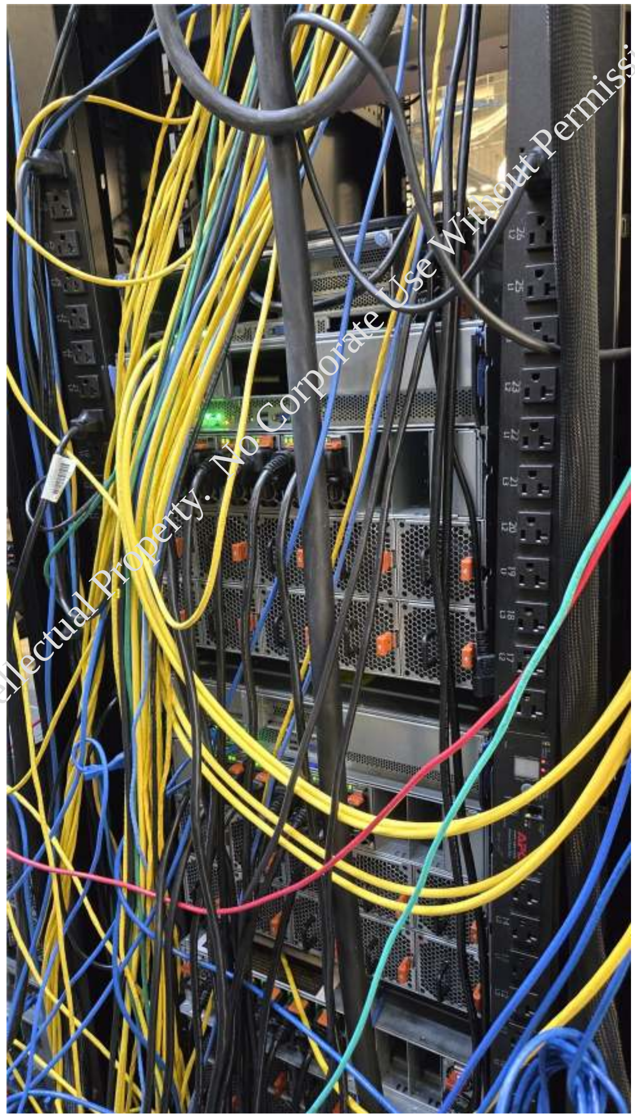

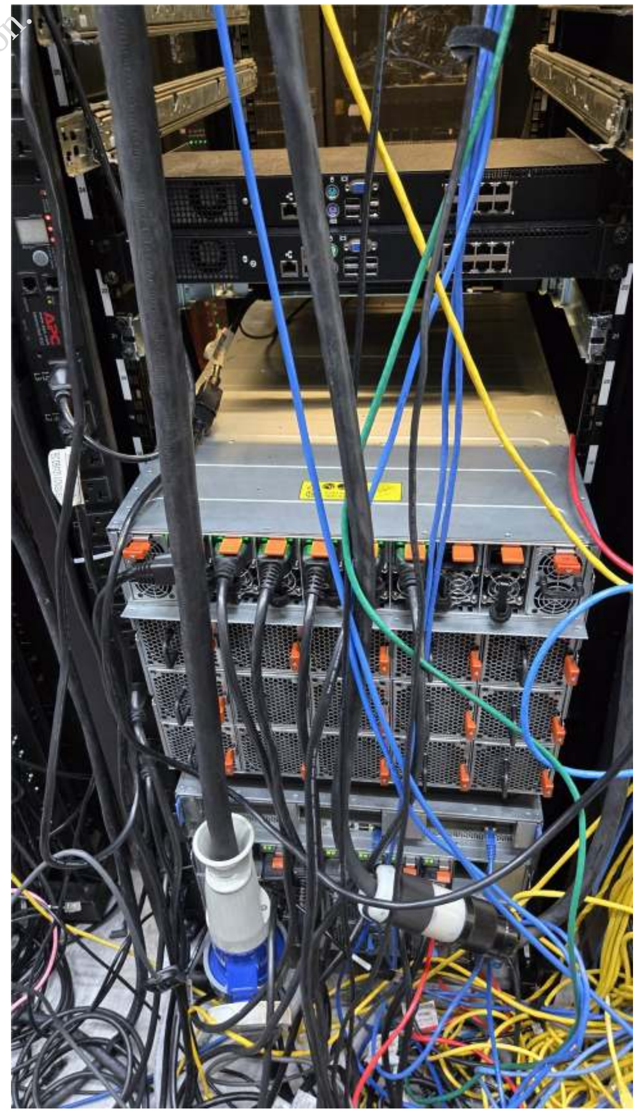

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

The trip hazards, blockage of limited aisle space, along with the hot spot zones, increased fan utilization of hot air, melting the cables blocking the airflow

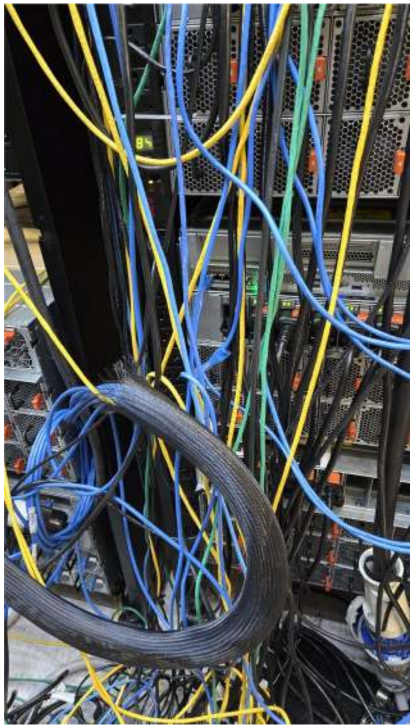

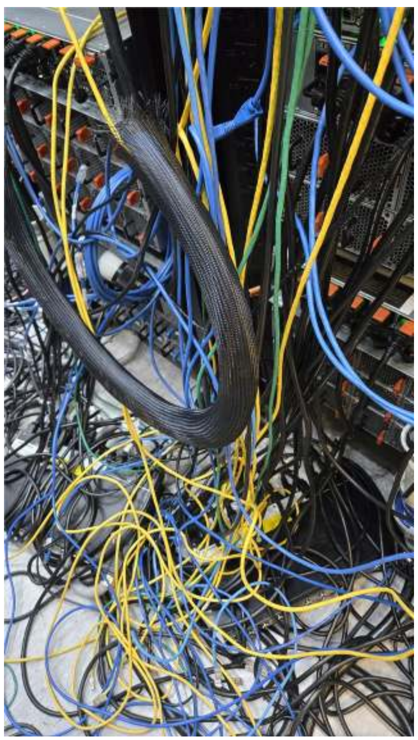

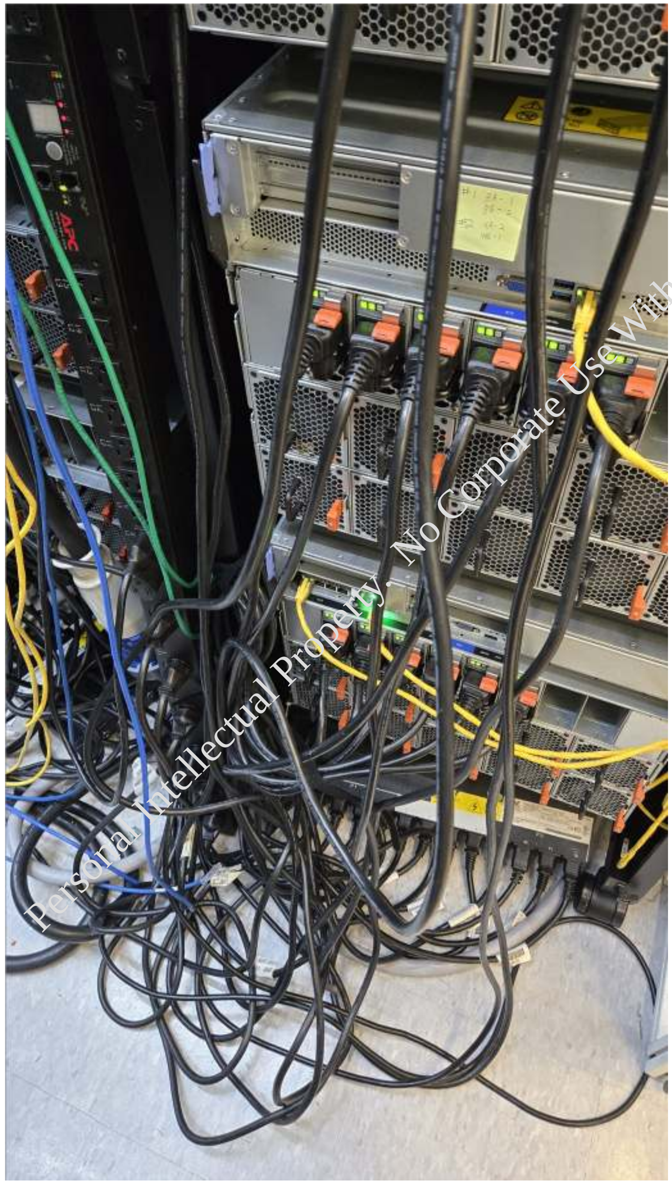

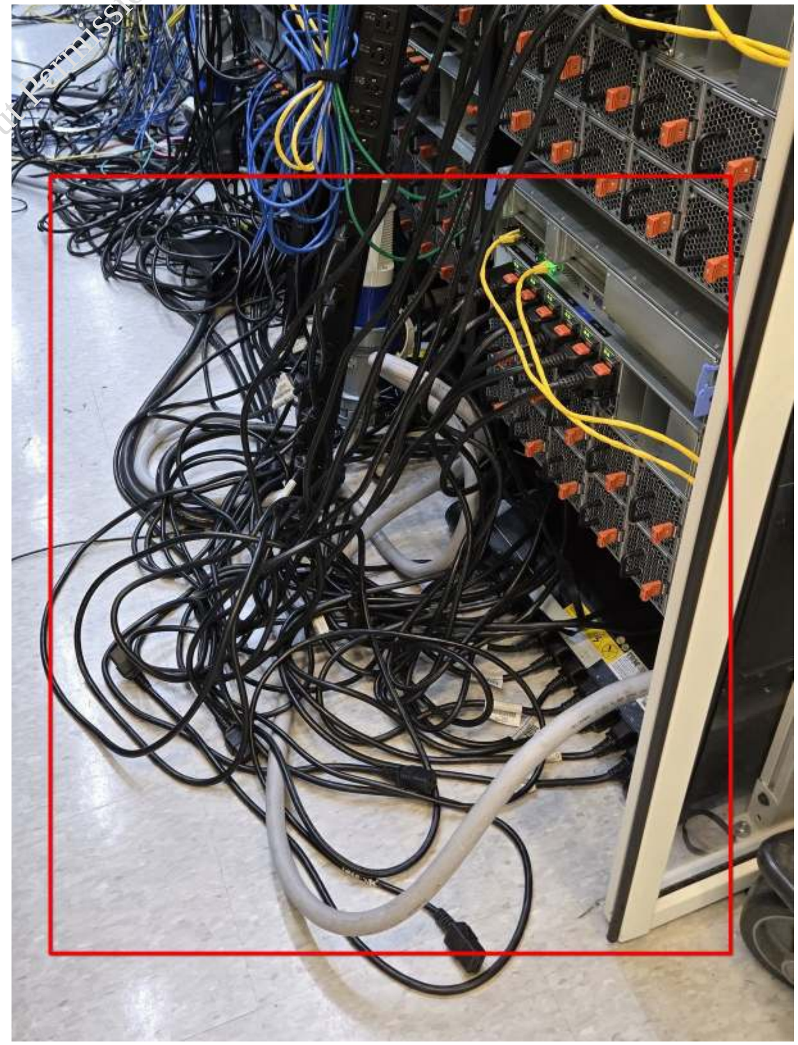

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

The system chassis / equipment is not in the proper rack size to accommodate the depth making any proper best practices cable management / air flow a best effort attempt given the constraints of the environment

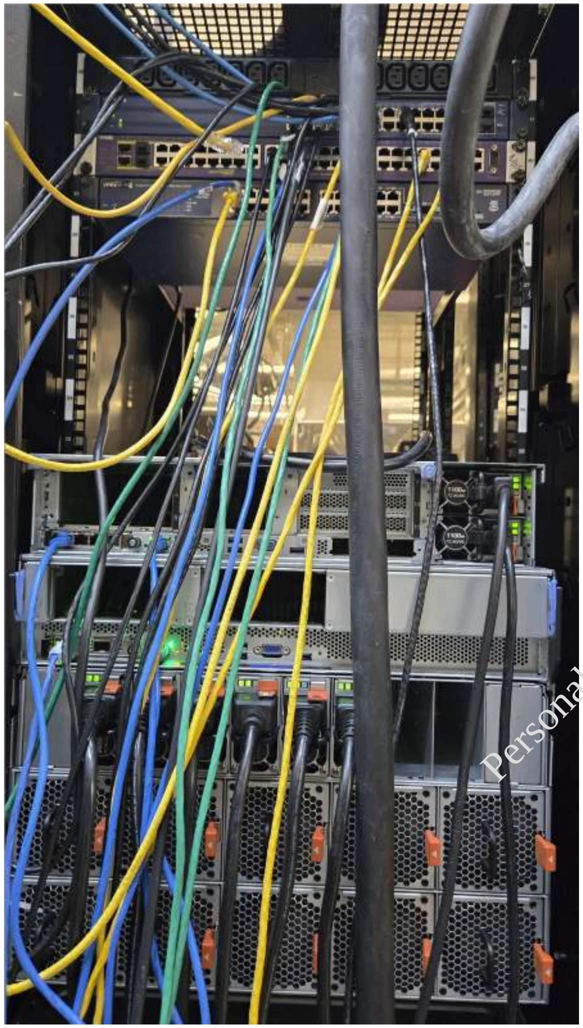

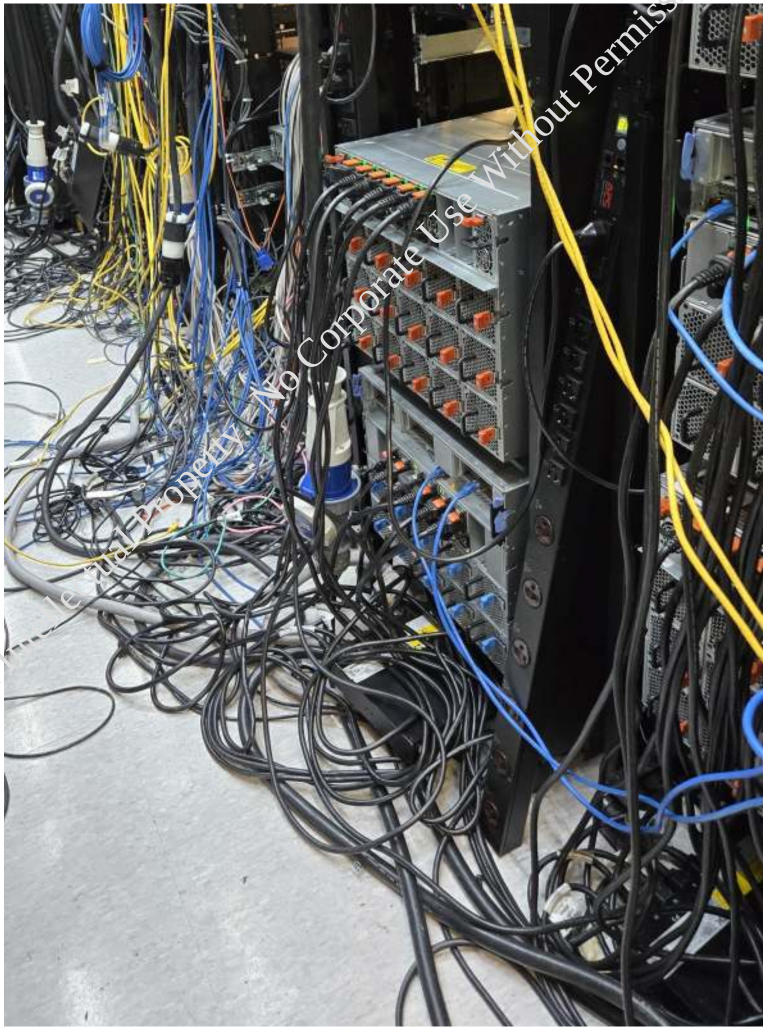

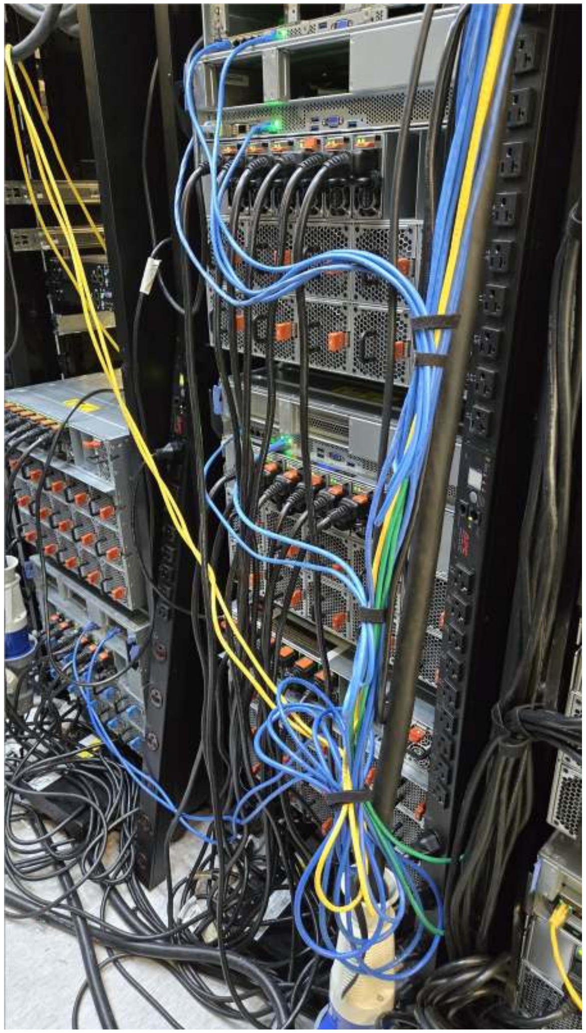

No Corporate Use Without Permission.

Working on remedying the network cables for some consistency on cable length, and switch port structure and population scheme

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

Redeployed physical cabling, decoupled inter-system connection dependencies, remedy the thermal air flow issues, given psu limitations, reclaimed limited aisle space

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

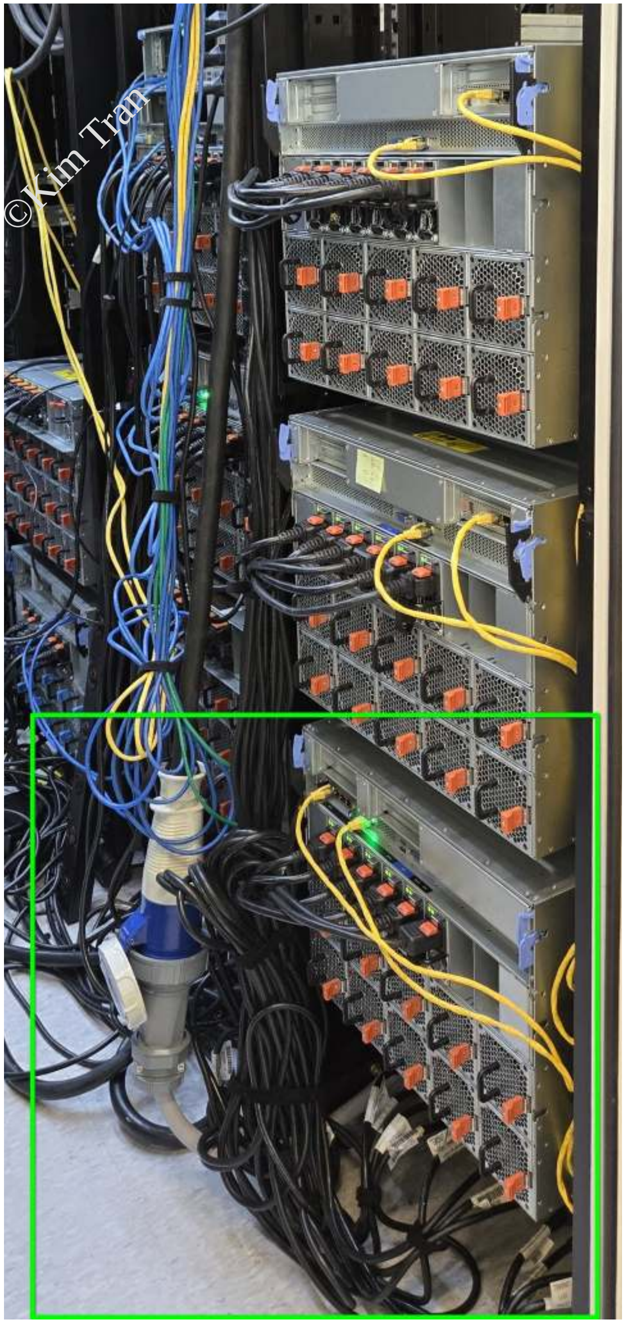

Redeployed cabling for racks 2 and 3, remedy the thermal air flow issues and reclaimed what little aisle space available

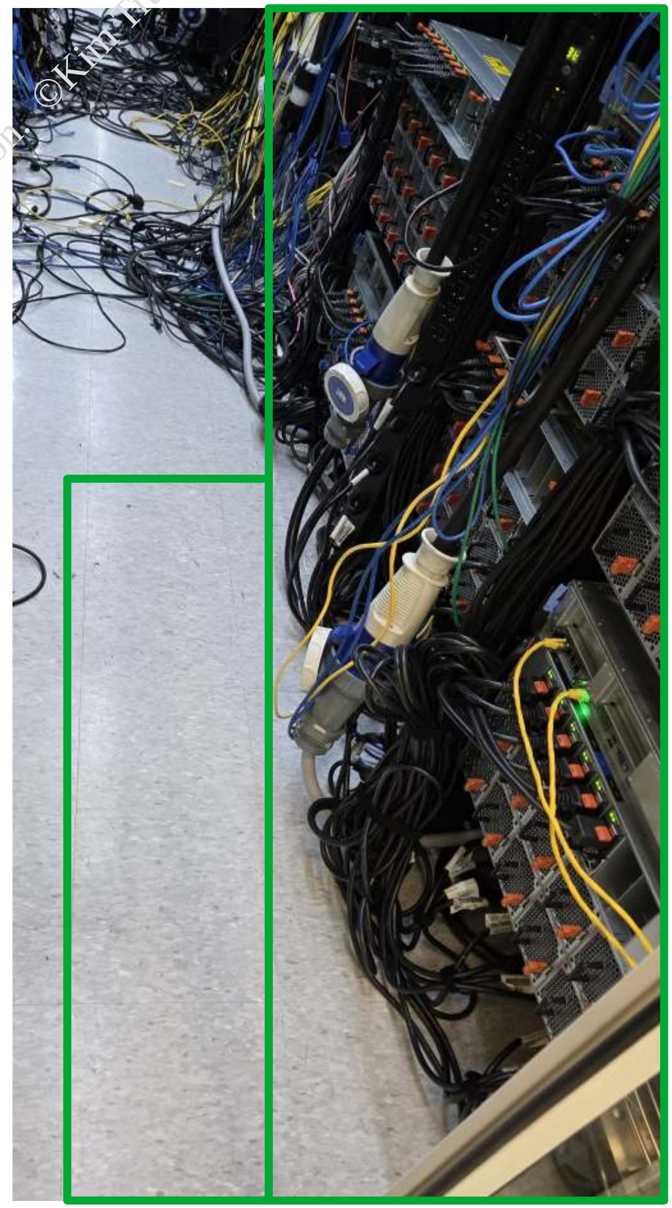

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran

Remedy the internal air flow / pressure as designed and engineered

Personal Intellectual Property. No Corporate Use Without Permission. ©Kim Tran